The Norton Introduction to

PHILOSOPHY

GIDEON ROSEN
Princeton University

ALEX BYRNE
Massachusetts Institute of Technology

JOSHUA COHEN
Apple University

SEANA VALENTINE SHIFFRIN
University of California, Los Angeles

W. W. NORTON & COMPANY, INC.
New York • London

W. W. Norton & Company has been independent since its founding in 1923, when William Warder Norton and Mary D. Herter Norton first published lectures delivered at the People's Institute, the adult education division of New York City's Cooper Union. The firm soon expanded its program beyond the Institute, publishing books by celebrated academics from America and abroad. By midcentury, the two major pillars of Norton's publishing program—trade books and college texts—were firmly established. In the 1950s, the Norton family transferred control of the company to its employees, and today—with a staff of four hundred and a comparable number of trade, college, and professional titles published each year—W. W. Norton & Company stands as the largest and oldest publishing house owned wholly by its employees.

Copyright © 2015 by W. W. Norton & Company, Inc.
All rights reserved
Printed in the United States of America
First Edition

Editor: Peter J. Simon
Associate Editor: Quynh Do
Project Editor: Diane Cipollone
Editorial Assistant: Gerra Goff
Managing Editor, College: Marian Johnson
Managing Editor, College Digital Media: Kim Yi
Production Manager: Vanessa Nuttry
Electronic Media Editor: Cliff Landesman
Marketing Manager, Philosophy: Michael Moss
Design Director: Hope Miller Goodell
Permissions Manager: Megan Jackson
Composition: S4Carlisle Publishing Services
Manufacturing: Quad/Graphics-Taunton

Permission to use copyrighted material is included on page 1125.

Library of Congress Cataloging-in-Publication Data
The Norton introduction to philosophy / Gideon Rosen, Princeton University; Alex Byrne, Massachusetts Institute of Technology; Joshua Cohen, Stanford University; Seana Valentine Shiffrin, University of California, Los Angeles.—First Edition.
pages cm
Includes bibliographical references and index.
ISBN 978-0-393-93220-1 (pbk. : alk. paper)
1. Philosophy—Textbooks. I. Rosen, Gideon A., editor. II. Byrne, Alex, 1960- editor. III. Cohen, Joshua, 1951- editor. IV. Shiffrin, Seana Valentine, editor.
BD21.N67 2015
100—dc23
2014046330

W. W. Norton & Company, Inc., 500 Fifth Avenue, New York, NY 10110
wwnorton.com

W. W. Norton & Company Ltd., Castle House, 75/76 Wells Street, London W1T 3QT

1 2 3 4 5 6 7 8 9 0

The Norton Introduction to

PHILOSOPHY

GIDEON ROSEN
Princeton University

ALEX BYRNE
Massachusetts Institute of Technology

JOSHUA COHEN
Apple University

SEANA VALENTINE SHIFFRIN
University of California, Los Angeles

2015

W. W. NORTON & COMPANY, INC.
New York • London

Harman, Elizabeth (2015), "Is it Reasonable to 'Rely on Intuitions in Ethics?", in Gideon Rosen, Alex Byrne, Joshua Cohen, and Seana Valentine Shiffrin, eds., *The Norton Introduction to Philosophy*, London: Norton & Company.

Elizabeth Harman (born 1975)

Harman is a Professor of Philosophy at Princeton University. She specializes in moral philosophy and is especially well known for her work on harm and the ethics of procreation.

IS IT REASONABLE TO "RELY ON INTUITIONS" IN ETHICS?

Some philosophers argue for ethical conclusions by relying on specific ethical claims about described cases. I will discuss and defend this practice. It is often described as "relying on intuitions," though I will argue that this description is deeply misleading.

Ethical arguments can usefully rely on specific ethical claims about described cases in at least three ways. First, a specific ethical claim can be offered as a *counterexample* to a more general ethical claim. Second, a specific ethical claim can

be used to motivate or *support* a more general ethical claim; for example, an author might argue that if a specific ethical claim is true, then it must be true *because* a more general ethical claim is true, and so the more general ethical claim must be true. Third, a specific ethical claim may be used in an argument for another specific ethical claim. There are many ways such an argument might proceed. For example, it might proceed via argument for a more general ethical claim, or it might proceed by claiming that there are no morally significant differences between the two cases in question that could warrant different verdicts about the cases.

It will be helpful to have in mind some examples of the type of argument I am discussing. In "Famine, Affluence, and Morality," Peter Singer argues that each of us ought to give a lot of money (much more than people typically give) to famine relief.[1] His argument relies on the claim that if a man is walking by a drowning child, and he is the only person in a position to save the child, but saving the child would involve getting his suit muddy, then he ought to save the child. This specific ethical claim is used to *support* the general claim that if one can prevent something bad from happening without sacrificing anything of comparable moral importance, then one should do so, which Singer then uses to support his conclusion about famine relief. Judith Jarvis Thomson's paper "A Defense of Abortion" argues for the claim that ordinary abortions are permissible even if early fetuses have the full moral status of persons.[2] Her argument crucially depends on the claim that if a man wakes up in a hospital, perfectly healthy himself but with his kidneys being used to keep a famous violinist alive, and if the violinist will die unless the man stays in the hospital for nine months, then it is permissible for the man to detach himself and leave the hospital, causing the violinist's death. She uses this claim as a *counterexample* to the general claim that it is never permissible to violate a right to life merely in order to have control of one's body. Both Singer's and Thomson's papers argue *from one specific ethical claim to another*.

Why might it be thought to be unreasonable to rely on specific claims about described cases in ethics? It is sometimes pointed out that not everyone agrees on these specific claims. Indeed, sometimes survey data is produced to *prove* that there is disagreement about the specific claims.[3] This complaint may simply misunderstand what is going on when people rely on specific claims about described cases in ethics, and in philosophy more generally. Such arguments have as a premise *a certain claim about the case*. They do not have as a premise *a claim that everyone agrees with a certain claim about the case*.

1. Peter Singer, "Famine, Affluence, and Morality," *Philosophy and Public Affairs* 1 (1972): 229–43. Excerpted in chapter 18 of this volume. [Harman's note.]

2. Judith Jarvis Thomson, "A Defense of Abortion," *Philosophy and Public Affairs* 1 (1971): 47–66. Excerpted in chapter 18 of this volume. [Harman's note.]

3. See Joshua D. Greene, "The Secret Joke of Kant's Soul," in *Moral Psychology*, vol. 3: *The Neuroscience of Morality: Emotion, Disease, and Development*, ed. W. Sinnott-Armstrong (MIT Press, 2007), which makes a more complicated argument than the one I go on to discuss; see Selim Berker, "The Normative Insignificance of Neuroscience," *Philosophy and Public Affairs* 37, no. 4 (2009): 293–329, for a critique of Greene. [Harman's note.]

It might be thought that philosophical arguments should not have any premises about which there is disagreement. But having such premises is in the very nature of philosophical arguments, and certainly ethical arguments. Many moral philosophy papers begin by assuming Kantianism, utilitarianism, consequentialism, or virtue ethics and proceeding from there. The arguments these papers offer are not bad arguments simply because they have deeply controversial assumptions. In ethics, and in philosophy more generally, there is quite a lot of disagreement. Arguments often have substantive premises with which some people agree and some people disagree. Two things happen when someone puts forward such an argument. Those who believe the premises are invited to follow the line of reasoning outlined, and to believe the conclusion of the argument. Those who do not believe the premises are invited to follow a related line of reasoning, to a weaker conclusion: they are invited to believe that *if* those premises are true, *then* the argument's conclusion is true. Thus, a paper that argues from particular premises is not interesting only to those who believe the paper's premises; it has something to say to *everyone*. And the weaker conclusion (that if the premises are true then the conclusion is true) may be interesting to someone even if he does not believe the argument's premises. For example, if one of the premises is (or is implied by) a view that he thinks is false, and that he wants to convince others is false, then the paper may help him in this project by showing further implausible commitments of the rival view. Or to return to our prior examples, Singer intends his 1972 paper to convince his readers of his conclusion. But some readers may become convinced of the weaker conditional: if it is morally obligatory to save a drowning child right in front of one, at the cost of getting one's suit muddy, then it is morally obligatory to give lots of money away regularly to prevent remote children from starving. These readers may believe this conditional and be moved to employ modus tollens rather than modus ponens: because we are not morally obligated to give lots of money away, they conclude that a man who can save a drowning child only at the cost of muddying his suit is not obligated to do so.

Similarly, someone who lacks a clear belief about whether it is permissible, in Thomson's violinist case, for the man in the hospital to unplug himself, may nevertheless find himself persuaded by Thomson's argument to believe that if the man's unplugging himself is permissible, then abortion of a fetus with full moral status is permissible. This person thereby comes to believe something about the relationship between the obligations to be a good Samaritan who helps others at significant cost to himself, on the one hand, and the obligation not to abort on the other: if the first obligation does not exist, then the second does not either.

Disagreement does pose another worry, which cannot be so easily dismissed. The fact that there is disagreement over a certain premise may seem to give those who believe the premise sufficient reason to doubt their judgment, so that they should cease to believe the premise.[4] This is an *epistemological worry*, that is, a worry about whether *belief* in the premise of an argument is justified; it is not a worry about

4. Jonathan Weinberg, Shaun Nichols, and Stephen Stich, "Normativity and Epistemic Intuitions," *Philosophical Topics* 29, no. 1/2 (2001): 429–60. [Harman's note.]

whether the premise is true. Epistemological worries pose the most serious kind of challenge to the practice I am defending. I now turn to three sorts of epistemological worries: the first arises from the fact of disagreement; the second arises from a concern that mere intuitive seemings cannot justify beliefs; and the third arises from a concern that some described cases are too far-fetched.

I will begin with the worry arising from disagreement.

Suppose that Anne believes a particular specific ethical claim, about the following case. A train is heading for five innocent people caught on the tracks, all of whom will die if they are hit. A person is standing on a bridge over the tracks. She can push a large fat man, who is next to her, off the bridge onto the track. His body would stop the train, preventing it from hitting the five. (Her own smaller body would not stop the train; if she jumped, she would die along with the five.) Anne believes it would be wrong for the woman to push the fat man off the bridge, to save the five. Suppose that Anne learns that there is disagreement about this case. In surveys, while many people agree with Anne, many disagree with her.[5] Furthermore, the disagreement is not just among people considering the case for the first time. Even among people who have thought long and hard about this case, there is disagreement.

Upon learning this, Anne may find herself in a situation in which apparent *epistemic peers* disagree with her. That is, Anne may expect that other ordinary people would be roughly as good as she is at discerning the moral truth about a described case if they have the same evidence she does; and she may take others to have the same evidence she has. It may seem that in this situation, Anne would be unreasonable in continuing to believe that it would be wrong to push the fat man. After all, she has no particular reason to think she is better at responding to the shared evidence than other people; but she would have to have such a reason to trust her own judgment more than others.

The view I have just outlined holds that when one faces disagreement from an epistemic peer, one should suspend judgment about the disputed claim, because one has no independent reason to take one's own judgment to be better. This view might be supported by the claim that when Anne confronts disagreement, the only thing that grounds her belief is *the fact that she has judged the claim to be true;* the question then seems to become whether she believes her judgment is better than those she disagrees with. This claim is not true, however. Independently of the fact of disagreement (and prior to Anne's learning of the disagreement), Anne's belief was either justified or not. If it was justified, there were some factors which made it justified. Those factors remain after Anne hears of the disagreement; what is under dispute is whether they are still sufficient to justify her belief, that is, whether they are undermined by the fact of disagreement. On the view I endorse, these factors *do* make it reasonable for Anne to continue to hold her belief; *they* furthermore justify her belief that others, in this case, are in error: she concludes they are in error because they think it is permissible to push the fat man, which is false. (The fact that she judged the claim to be true plays no justifying role in her continuing to believe the claim.) On the other hand, if Anne's belief was not initially justified, then it is still not

5. Greene, 2007. [Harman's note.]

justified after she learns of the disagreement. But the fact of disagreement does not make it unjustified; it was independently unjustified.

I have endorsed a stark view on which disagreement poses no skeptical threat at all.[6] If one's beliefs are justified, learning that some others disagree should not at all undermine one's beliefs. I might instead have endorsed a more concessive view, according to which the fact that some others disagree with one makes it reasonable to be *less certain* of a belief one holds, but does not require one to suspend belief.[7] Both this more concessive view and the stark view I favor vindicate Anne in continuing to believe it is wrong to push the fat man, even in the face of disagreement.[8]

Let's turn to the second epistemological worry. This worry arises out of two claims. First, all specific ethical beliefs are formed on the basis of an intuitive seeming — on the basis of a claim's seeming to be true, but not for any other reason, that is, not on the basis of any evidence other than the claim's seeming to be true. Second, it is unreasonable to form specific ethical beliefs on the basis of intuitive seemings.

Is it true that all specific ethical beliefs are formed on the basis of intuitive seemings? No. Some specific ethical beliefs are formed on the basis of explicit reasoning *from ethical theories,* despite the believer's finding their negations intuitive. For example, a consequentialist may believe it is permissible to push the fat man off the bridge despite finding it intuitive that it is impermissible.

We might revise the worry's first claim to this: if a specific ethical belief is one that a person simply finds himself with upon reading a description of a case, then it was formed on the basis of an intuitive seeming — on the basis of the claim's seeming to be true, but not for any other reason. The worry's second claim remains the same, that it is unreasonable to form specific ethical beliefs on the basis of intuitive seemings. The worry's conclusion becomes the more limited claim that *in cases where a belief was formed in this way,* the belief is not a reasonable belief and so not a reasonable basis for an inference to a new belief. (The worry in this form applies to an argument as read by some readers but not as read by others.)

Is it true that if a reader simply finds himself with an ethical belief about a case upon reading a description of the case, then the belief must have been formed on the basis of an intuitive seeming — on the basis of the claim's seeming to be true, but not for any other reason? Surely not. The belief may have been implicitly inferred from other ethical commitments the person has;[9] this may have happened without the person's realizing it, or he may be *unsure* whether this has happened.

No *general* account will accurately describe what happens in all instances in which people read descriptions of cases and then find themselves with beliefs about the cases. Sometimes people have pre-existing beliefs about the cases. (For example,

6. Thomas Kelly, "The Epistemic Significance of Disagreement," in *Oxford Studies in Epistemology,* vol. 1, ed. John Hawthorne and Tamar Gendler Szabo (Oxford University Press, 2005), 167–96. [Harman's note.]

7. Thomas Kelly, "Peer Disagreement and Higher Order Evidence," in *Disagreement,* ed. Richard Feldman and Ted Warfield (Oxford University Press, 2010). [Harman's note.]

8. For a defense of the view that Anne would not be reasonable to continue to hold her belief, see Adam Elga, "Reflection and Disagreement." *Noûs* 41, no. 3 (2007): 478–502. [Harman's note.]

9. F. M. Kamm, "Introduction," in *Creation and Abortion* (Oxford University Press, 1992). [Harman's note.]

some people already believe that a stranger walking by a drowning child should save the child, before reading Singer's description of that case.) Sometimes people infer particular beliefs about the cases from more general beliefs they already have. Or they make explicit a particular belief they already implicitly had.

Consider the famous example of the doctor who cuts up an innocent healthy person to save five people dying of organ failure. It regularly happens that people hear this case described for the first time. They often react by believing that what the doctor does is wrong. Though they have never heard the case described, they do not *newly believe* that what the doctor does is wrong; they either already implicitly believed it or were already implicitly committed to it.

These points show that someone may simply find herself with a particular belief, upon hearing a case described, without it being true that her belief was formed on the basis of an intuitive seeming — on the basis of the claim's seeming true to her, and nothing else. We should not assume that, most of the time, beliefs about described cases are formed on the basis of intuitive seemings.

Indeed, intuitive seemings may play *no role* in the epistemology of specific ethical beliefs, as they actually occur. But if intuitive seemings do play a role, this may not be problematic. We might have the view that intuitive seemings *can* justify beliefs. We might hold that intuitive seemings justify beliefs in the same way that perceptual seemings justify beliefs. This claim could be elaborated in several different ways. One view holds that a person's intuitive seemings justify beliefs if her intuitive seemings are reliable;[1] this is an analogue of a reliabilist story about perceptual justification. Another view would hold that intuitive seemings justify beliefs simply because of their content: that it seems to you that p is true is in itself a reason to believe that p is true; this is an analogue of the "dogmatic" view of perceptual justification offered by James Pryor.[2]

We are now in a position to see why the practice I am defending — the making of ethical arguments that rely on specific ethical claims about described cases — should not be described as "relying on intuitions" in doing ethics. At least two mistakes are present in that description. First — as I mentioned early in this paper — there is a conflation of two very different practices: relying on certain claims that may in fact be intuitive, on the one hand, and relying on the claim *that* these claims are intuitive, on the other hand. Philosophical arguments of the type I am defending (and of the type often criticized for "relying on intuitions") do not rely on any claims about intuition; they rely on specific moral claims themselves. Second, there is an assumption that whenever we believe specific moral claims about described cases, we believe them simply because they are intuitive; as I've just argued, this is not true.

The third epistemological worry is sometimes voiced as follows: "Some described cases are too odd, too complicated, or involve too much science fiction for us to have reasonable beliefs about the cases." If this worry is meant to apply to all uses of claims about described cases, it is false. Singer's case, of the man and the drowning child, is neither odd, nor complicated, nor involving science fiction. Thomson's case

1. George Bealer, "A Theory of the A Priori," *Philosophical Perspectives* 13 (1999): 29–55. [Harman's note.]

2. James Pryor, "The Skeptic and the Dogmatist," *Noûs* 34, no. 4 (2000): 517–49. [Harman's note.]

involves some *fiction*: she supposes that medical records have shown that the man in the hospital, attached to the violinist, is the only person who can help; she supposes that he has been kidnapped and attached; and she supposes that one person's kidneys can be used to restore another person to health across nine months. But none of these suppositions is very far removed from the actual world: donor databases for bone marrow transplants sometimes reveal to someone that he is the only person who can save a stranger's life, organs are sometimes stolen from healthy people, and in order to donate kidneys or parts of livers, people sometimes endure serious health risks and hospital stays to restore others to health. The case of the woman on the bridge with the fat man is not very complicated. The case of the doctor who cuts up his healthy patient to save five patients with organ failure is also quite simple. Neither of these two cases involves any science fiction.

Sometimes people complain that particular described cases are too odd, too complicated, or involve too much science fiction as a way of explaining why they themselves are unmoved by the arguments that rely on claims about these cases: they find that they *lack* beliefs about the cases, or that they lack stable, confident beliefs about the cases. This fact, that *some people* lack beliefs about the claims in question, does not show that everyone lacks such beliefs. The arguments may nevertheless be perfectly good arguments that appeal to claims believed by some people but not by everyone.

It is a *much stronger* claim that some particular described cases are so odd, so complicated, or involve so much science fiction that *no one* could reasonably have a justified belief about the cases. This claim is not true about any of the four described cases I have just mentioned. Indeed, each case is such that we could come upon a similar case in real life. Would the case be too odd for us to judge whether the agent had acted permissibly? Surely not.[3]

I will briefly mention a final concern. This concern maintains that our general ethical beliefs are better grounded, or more reliable, than our specific ethical beliefs about cases, such that we should always do ethics by proceeding *from* the general to the specific, and never vice versa.[4] Kantians and utilitarians often believe this is the correct view of the epistemology of ethics. There are many different reasons that might be offered for this view. But none of the three epistemological worries I have discussed can support this view. As for the first worry, I conjecture that if we were to conduct surveys of people's beliefs in general ethical claims, we would find substantial disagreement. Consideration of disagreement will not tell in favor of general ethical claims over specific ethical claims. As for the second worry, whether intuitive seemings can be a source of justified beliefs favors neither specific nor general ethical claims. The third worry might seem to favor general claims. Whereas some specific ethical claims are about cases that may be odd or complicated, nothing corresponding holds of general ethical claims: being general,

3. Timothy Williamson, "Thought Experiments," chapter 6 in *The Philosophy of Philosophy* (Blackwell, 2007). [Harman's note.]

4. Peter Singer, "Ethics and Intuitions," *The Journal of Ethics* 9 (2005): 331–52; R. M. Hare, "Rawls' *A Theory of Justice*: Part I," *Philosophical Quarterly* 23 (1973): 144–55. [Harman's note.]

they are not focused on anything odd or complicated. But the third worry, properly understood, simply presses the point that some specific ethical claims are such that some people will be unsure what to make of them. The same is true of general ethical claims.

Conclusion

I have defended the making of arguments for ethical conclusions on the basis of specific ethical claims about described cases. I have argued that three objections to this practice fail.[5]

Where does this leave someone who wants to pursue ethical questions by reading philosophy papers? Suppose you read a paper that argues for an ethical conclusion on the basis of specific ethical claims about described cases.

If you believe the specific ethical claims the paper relies on, then the paper is offering you an argument for its conclusion, which you might well reasonably rely on in coming to believe the conclusion.

If you do not believe a specific ethical claim the paper relies on, either because you believe it is false or because you are unsure what to make of it, then the paper is not in a position to convince you of its conclusion. You may, however, find it interesting whether the weaker conditional is true — that if its premises are true, then its conclusion is true — and the paper may reasonably convince you of this claim. You may also find it interesting that this paper may convince others; you may want to engage with the paper as an argument addressed to others, perhaps showing (to those who believe the premises) that the premises do not really imply the conclusion.

When is a philosophy paper that relies on specific ethical claims criticizable for relying on such claims? No paper is criticizable simply for being *a paper that relies on some specific ethical claims about described cases* — some papers do so and are excellent, compelling, important philosophy papers.[6] If a paper relies on claims that are only believed by some people, it is still a contribution for being of interest to those people (and it is also of interest to everyone as offering a claim about what follows from certain other claims). But a paper might rely on specific ethical claims such that no one is in a position to form a justified belief about those claims. Such a paper is less interesting because it provides no interesting argument for its conclusion; but it may still be interesting as arguing that if its premises are true, then its conclusion is true.

5. For two very different defenses of reliance of specific ethical claims, see Kamm, 1992 and the final section of Tamar Szabo Gendler, "Philosophical Thought Experiments, Intuitions, and Cognitive Equilibrium," *Midwest Studies in Philosophy* 31 (2007): 68–89. [Harman's note.]

6. Thomson, 1971; Singer, 1972; Michael Tooley, "Abortion and Infanticide," *Philosophy and Public Affairs* 2 (1972): 37–65; James Rachels, "Active and Passive Euthanasia," *New England Journal of Medicine* 292 (1975): 78–80; Susan Wolf, "Asymmetrical Freedom," *Journal of Philosophy* 77, no. 3 (1980): 151–166; Derek Parfit, "Part Three: Personal Identity," *Reasons and Persons* (Oxford University Press, 1982); Seana Shiffrin, "Wrongful Life, Procreative Responsibility, and the Significance of Harm," *Legal Theory* 5 (1999): 117–48; and Frances Kamm, "Terrorism and Several Moral Distinctions," *Legal Theory* 12, no. 1 (2006): 19–69. [Harman's note.]